

Distributed Systems

Events

def An event is an occurrence at a single process that causes a transition in its local state. It is formally represented as:

$$e_i^k = \text{k-th event of process } i$$

Relationship "Happened Before"


def The event e_i^l "happened before" e_j^m if:

① $i = j$ and $l < m$ (same process and e_i^l happens before e_i^m)

OR

② e_i^l is the sender of the message received by e_j^m

is this def complete? Based on that $e_2^1 \rightarrow e_1^1$



Local State

def The Local State of a process is simply the "snapshot" of the process in a specific situation. It is formally represented as:


$$\sigma_i^k = \text{state of process } i \text{ after } k \text{ events } (e_1^1, \dots, e_i^k)$$

Global State

def The Global State is a "snapshot" of the entire system. It is constructed from the union of the local state of every process. It is formally represented as:

$$\sigma^k = \text{state of all processes after } k \text{ events} = (\sigma_1^k, \sigma_2^k, \dots, \sigma_n^k)$$


N.B. A global state is actually a "straight cut" on the time



$$\sigma^4 = (\sigma_1^4, \sigma_2^4, \sigma_3^4)$$

Cut

def A "cut" is a set of events that contain the first k_i events of process i for each process



$$C = (\sigma_1^5, \sigma_2^4, \sigma_3^4)$$

However, we are not interested in "every cut". We only need "consistent cuts".

Consistent Cuts


def A cut C is consistent iff: $e \xrightarrow{\text{happens before}} e' \wedge e' \in C \Rightarrow e \in C$

if and only if *happens before*


if e' receives a message from e , both e and e' must be in the cut

NON necessariamente il contrario. Se e manda un messaggio a e' ed $e \in C$, e' potrebbe anche NON appartenere a C

✓ CONSISTENT



✗ NOT CONSISTENT



Run

def A "run" is a sequence of events that represents a specific execution

As before, we can also define "what is a consistent run?"

def A run R is consistent iff: $e \xrightarrow{\text{happens before}} e' \Rightarrow e \text{ appears before } e' \text{ in } R$

if and only if *happens before*

And, basically, a Run is a sequence of Cuts. So, we could cut a Run to get a Cut.


$$R = \overbrace{e_1^1 \ e_1^2 \ e_1^3}^{C'} \mid e_1^4 \ e_1^5 \ e_1^6 \dots$$

Lamport Clock

Let's start from an idealized world where capturing a consistent snapshot of the system's global state is very easy, thanks to 2 different assumptions:

- GLOBAL CLOCK: All processes share the exact same physical clock, meaning the clock drift is negligible or zero ($\Delta d = 0$)
- SYNCHRONOUS NETWORK: There is a known maximum Δt time needed to deliver messages.

Under these assumptions the following simple protocol works properly:




- 1) A coordinator process (p_0) broadcasts a request at time t_0 , saying: "Send me your local state at time $t = t_0 + \Delta t$ "
- 2) Because the network is synchronous all processes receive the request in time (before Δt time)
- 3) Thanks to the global clock all processes record their local state at the exact same time

However, in real world distributed systems this protocol cannot actually work. Physical clocks inevitably drifts and network congestion, failures, retransmissions etc... make the message delays unpredictable. So, we need to change the protocol in order to eliminate these 2 assumption.

Removing the Global Clock

IDEA: Instead of asking processes to record their state at a specific time, every process sends a message to p_0 whenever local event happens. Then, p_0 will reconstruct the run to create the global snapshot.



NOTE: At this moment we are still relying on the SYNCHRONOUS NETWORK assumption. Because message delay is bounded by Δt p_0 knows that any event older than Δt has definitely arrived.

Removing the Synchronous Assumption


IDEA: To understand the system's global state we do not actually need to know the exact timestamp an event occurred. We only care about the order of the events. So, we can introduce the Clock Condition and the Lamport Clock to reorder messages to their "logical clock".

- ^(WEAK) CLOCK CONDITION: $e \xrightarrow{\text{happened before}} e' \Rightarrow T(e) < T(e')$ (NOTE: $\Leftarrow$ it's not true)
- LAMPORTS CLOCK: Each process maintains an integer counter C and attaches it in every message.

- At each LOCAL event increment $\Rightarrow C = C + 1$

- Receiving a message with timestamp $t \Rightarrow C = \max(C, t) + 1$

Finally, to reconstruct the order of the events a node must ensure that no message with an older timestamp will arrive. To make this we can use a LAST MESSAGE VECTOR. Each process tracks the latest timestamp received from the others in a vector $[last_1, \dots, last_n]$. Events with timestamps up to $\min(last_1, \dots, last_n)$ can be safely committed because no future message can have smaller timestamp.



PROBLEM: If a process become idle its local counter C stops increasing. Consequently, its entry in the LAST MESSAGE VECTOR remains low freezing the global minimum. In that case the other processes can not commit their events because they are waiting on the silent process.

SOLUTION: Heartbeats / Periodic Null Message

(Strong) Clock Condition


The previous (weak) clock condition " $e \rightarrow e' \Rightarrow T(e) < T(e')$ " could not say if e happened before e' or if they were simply concurrent systems. To solve this we want to achieve a STRONG CLOCK CONDITION:

$$e \rightarrow e' \Leftrightarrow T_s(e) \leq T_s(e')$$

To achieve this bidirectional relationship we need to look at the entire history of an event.

Exercise: Is it possible that $e_i \not\rightarrow e_j \wedge e_j \not\rightarrow e_i$?

Yes!




Event's History

def The History of the event e is the set of events that happened before e . It is formally represented as:

$$H(e) = \{e' \text{ t.c. } e' \xrightarrow{\text{happened before}} e\}$$

NB if $e_i \rightarrow e_j$ then $e_i \in H(e_j) \wedge H(e_i) \subseteq H(e_j)$


Vector Clock



Now we can define $Ts(e_z^8) = H(e_z^8) = \{e_1^5, e_2^8\}$ and write it "compressed" as a vector (known as Vector Clock) where each element v_i represents the latest known logical clock of the process p_i :

$$Ts(e_z^8) = H(e_z^8) = \begin{bmatrix} 5 \\ 8 \end{bmatrix} \begin{matrix} \rightarrow p_1 \\ \rightarrow p_2 \end{matrix}$$


Exercise: Compute the History of each event



At this point we can also define that: "A vector is smaller than or equal to another iff every single component is smaller then or equal to its corresponding component"
Formally:

$$\begin{bmatrix} x \\ y \end{bmatrix} \leq \begin{bmatrix} x' \\ y' \end{bmatrix} \Leftrightarrow x \leq x' \wedge y \leq y'$$

Because Vector Clock perfectly track all the order dependences, the coordinator p_0 will be able to detect if messages have arrived out of order (and potentially reorder them).



In this way p_0 can "reconstruct the Run" that, thanks to Vector Clock, it's also CONSISTENT!


Chandy - Lamport Clock

The Chandy - Lamport Clock allows to take a picture of the global state without halting the underlying execution.


IDEA:

- 1) Each process sends the "take a snapshot" message to everybody else AS SOON AS it receives the message

↳ Example: p_2 sends the message as soon as it is received by p_1 , and not from p_0



2) The cut on the exact moment each process take the snapshot forms the global state C .




But: "is this cut C consistent?"


Dim. To prove that the Chandy-Lamport algorithm always generates a consistent cut we can use a proof by contradiction:

$$\text{HP: } e \rightarrow e' \wedge e' \in C \Rightarrow e \in C$$


Since we are assuming that $e \notin C$, the cut C must have been taken **BEFORE** e .



On the contrary, since we know $e' \in C$, the cut C must have been taken **AFTER** e' .



However, the Chandy-Lamport Clock strongly relies on the assumption that communication channels are FIFO. So, because the snapshot marker is sent as soon as possible, if the snapshot was taken **BEFORE** e the message was also sent **before** e and, since the channel is FIFO and $e \rightarrow e'$, it must have arrived before e' .




So, we have proved that: $e \rightarrow e' \wedge e' \in C \Rightarrow e \in C$

So, the Chandy-Lamport Clock guarantees perfectly consistent cuts/global states.

Atomic Commit

Let's start from an example: The 2020 elections in Sapienza. During the covid pandemic these elections were done remotely. Consequently they needed a distributed system to collect the (vote, voter) pairs without storing the 2 information together (to ensure the vote anonymity).




IDEA:

- 1) register the voter on A
- 2) register the vote on B
- 3) delete the pair (vote, voter) on C

All of these operations SHOULD be done together, as a single ATOMIC COMMIT.
However A, B, C could fail during the Transaction... we need a protocol to ensure the atomic commit.

PROTOCOL:



However, we are considering a REALISTIC system (asynchronous and fair-loss). So, to make it work properly we need 4 different ASSUMPTIONS.

ASSUMPTIONS:

- SAFETY** AC1: All processes that reach a decision must reach the same one (ABORT or COMMIT)
- SAFETY** AC2: Once a process reach a decision it doesn't change
- SAFETY** AC3: COMMIT can be decided only if all processes votes YES
- LIVENESS** AC4: if there are NO FAILURES and everybody votes YES, the decision must be COMMIT

NOTE: if it was an "if and only if" the problem would not be solvable

FAILURES: But, again, we have to consider failures:

- **VOTE REQUEST DOESN'T ARRIVE:** We can simply set a TIMEOUT. If nothing arrives then Abort.
- **VOTE DOESN'T ARRIVE:** Also the coordinator can set a TIMEOUT. If nothing arrives then Abort.
- **DECISION DOESN'T ARRIVE:** In this case we cannot set the TIMEOUT... What if we ABORT and then the COMMIT DECISION arrives? All the other processes will commit and our part will be lost invalidating the system. We could proceed in different ways.
 - **COOPERATE:** thanks to AC1, we can ask to other processes. If anyone aborted we abort; if anyone committed we commit.
 - **WAIT FOREVER:** Since the channel is fair, sooner or later, the decision will arrive.
 - **IMPLEMENT A DISTRIBUTED LOG:** So that, in case of crash, we can reload the log and see what happened until now.

Distributed Log

Unfortunately, we cannot actually create a Distributed Log... It is like an abstraction derived from the collection of the Local Logs of each process. Note that each process (both coordinator or not) must save the local log on its permanent disk, to survive failures. Only by combining the local logs we can reconstruct the whole transaction.

- **LOG THEN SEND:** Generally, this is the safest approach. If there is a crash the process reads the log, sees that it was supposed to send a message and re-sends it.
- **SEND THEN LOG:** This is more risky. If there is a crash immediately after sending a message but before writing the log, after the rebooting the process will have no idea it performed that action.



Let's how this applies to the previous protocol:

- **PROCESSES:**
 - **IF THE VOTE WAS "YES":** It's mandatory the Log then Send. Otherwise, with the Send then Log, if the process votes YES, then fails and forget, could resend a different vote ("NO") invalidating the previous message at this point already received by the coordinator
 - **IF THE VOTE WAS "NO":** Log then Send is better but both are fine. With Send then Log if a process votes "NO" then fails and forget, after the reboot it will abort by default. So the vote doesn't change.
- **COORDINATOR:** The coordinator crashes, then reboot and read the log:
 - 1) **THERE ARE NO VOTES:** probably failed at ①. Re-sends the VOTE REQUEST
 - 2) **THERE IS AT LEAST ONE "NO":** Abort directly

- 3) THERE ARE ALL "YES" BUT SOME VOTE IS MISSING: We could have lost some "no" during the crash (or the vote was "yes" but will never arrive). The coordinator Abort for safety.
- 4) THERE ARE ALL "YES" BUT NO DECISION LOG: With Log then Send, no decision log mean no decision sent. So we could send COMMIT or ABORT. To stay safe lets assume that we ABORT always in this case. A Send then Log, instead, would be a disaster. Imagine that the coordinator sends "COMMIT" and then crash. After the reboot there is no log, it doesn't know the previous decision and sends "ABORT" for safety. The processes would receive both "COMMIT" and "ABORT"...
- 5) THERE ARE ALL "YES" AND THE DECISION LOG: In both cases, Log then Send and Send then Log, we simply retransmit the same DECISION as before.